

МОСКОВСКИЙ ГОСУДАРСТВЕННЫЙ ТЕХНИЧЕСКИЙ УНИВЕРСИТЕТ
им. Н.Э. Баумана

Факультет «Информатика и система управления»
Кафедра ИУ-5 «Системы обработки информации и управления»

ОТЧЕТ

Лабораторная работа №2
по курсу «Методы машинного обучения»

Исполнитель - студент группы ИУ5-21М:

Кауров Максим _____

Москва – 2020 год

▼ Часть 1

```
import numpy as np
import pandas as pd
import time
pd.set_option('display.max.columns', 100)
# to draw pictures in jupyter notebook
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
# we don't like warnings
# you can comment the following 2 lines if you'd like to
import warnings
warnings.filterwarnings('ignore')
```

```
#Загружаем датасет
from google.colab import files
files.upload()
```



Выбрать файлы adult.data.csv

- **adult.data.csv**(application/vnd.ms-excel) - 3518607 bytes, last modified: 06.04.2020 - 100% done
Saving adult.data.csv to adult.data.csv
{'adult.data.csv': b'age,workclass,fnlwgt,education,education-num,marital-status,occu

```
data = pd.read_csv('adult.data.csv')
data.head()
```



	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband
							

```
print("Список колонок с типами данных")
print(data.dtypes)
```



Список колонок с типами данных

```
age          int64
workclass    object
fnlwgt       int64
education    object
education-num int64
marital-status object
occupation   object
relationship object
race         object
sex          object
capital-gain  int64
capital-loss  int64
hours-per-week int64
native-country object
salary       object
dtype: object
```

1. How many men and women (sex feature) are represented in this dataset?

```
data['sex'].value_counts()
```

```
Male      21790
Female    10771
Name: sex, dtype: int64
```

2. What is the average age (age feature) of women?

```
data.loc[data['sex'] == 'Female', 'age'].mean()
```

```
36.85823043357163
```

3. What is the percentage of German citizens (native-country feature)?

```
float((data['native-country'] == 'Germany').sum()) / data.shape[0]
```

```
0.004207487485028101
```

4-5. What are the mean and standard deviation of age for those who earn more than 50K per year and those who earn less than 50K per year?

```
ages1 = data.loc[data['salary'] == '>50K', 'age']
ages2 = data.loc[data['salary'] == '<=50K', 'age']
print("The average age of the rich: {0} +- {1} years, poor - {2} +- {3} years.".format(
    round(ages1.mean()), round(ages1.std(), 1),
    round(ages2.mean()), round(ages2.std(), 1)))
```

```
↳
```

6. Is it true that people who earn more than 50K have at least high school education? (education: Assoc-voc, Masters or Doctorate feature)

```
data.loc[data['salary'] == '>50K', 'education'].unique() # Ответ: Нет
```

```
↳ array(['HS-grad', 'Masters', 'Bachelors', 'Some-college', 'Assoc-voc',  
        'Doctorate', 'Prof-school', 'Assoc-acdm', '7th-8th', '12th',  
        '10th', '11th', '9th', '5th-6th', '1st-4th'], dtype=object)
```

7. Display age statistics for each race (race feature) and each gender (sex feature). Use groupby method for men of Amer-Indian-Eskimo race.

```
for (race, sex), sub_df in data.groupby(['race', 'sex']):  
    print("Race: {0}, sex: {1}".format(race, sex))  
    print(sub_df['age'].describe())
```

```
↳
```

Race: Amer-Indian-Eskimo, sex: Female

count 119.000000
mean 37.117647
std 13.114991
min 17.000000
25% 27.000000
50% 36.000000
75% 46.000000
max 80.000000

Name: age, dtype: float64

Race: Amer-Indian-Eskimo, sex: Male

count 192.000000
mean 37.208333
std 12.049563
min 17.000000
25% 28.000000
50% 35.000000
75% 45.000000
max 82.000000

Name: age, dtype: float64

Race: Asian-Pac-Islander, sex: Female

count 346.000000
mean 35.089595
std 12.300845
min 17.000000
25% 25.000000
50% 33.000000
75% 43.750000
max 75.000000

Name: age, dtype: float64

Race: Asian-Pac-Islander, sex: Male

count 693.000000
mean 39.073593
std 12.883944
min 18.000000
25% 29.000000
50% 37.000000
75% 46.000000
max 90.000000

Name: age, dtype: float64

Race: Black, sex: Female

count 1555.000000
mean 37.854019
std 12.637197
min 17.000000
25% 28.000000
50% 37.000000
75% 46.000000
max 90.000000

Name: age, dtype: float64

Race: Black, sex: Male

count 1569.000000
mean 37.682600
std 12.882612
min 17.000000
25% 27.000000
50% 36.000000
75% 46.000000
max 90.000000

Name: age, dtype: float64

Race: Other, sex: Female

```

count    109.000000
mean      31.678899
std       11.631599
min       17.000000
25%       23.000000
50%       29.000000
75%       39.000000
max       74.000000
Name: age, dtype: float64
Race: Other, sex: Male
count    162.000000
mean      34.654321
std       11.355531
min       17.000000
25%       26.000000
50%       32.000000
75%       42.000000
max       77.000000
Name: age, dtype: float64
Race: White, sex: Female
count    8642.000000
mean      36.811618
std       14.329093
min       17.000000
25%       25.000000
50%       35.000000
75%       46.000000
max       90.000000
Name: age, dtype: float64
Race: White, sex: Male
count    19174.000000
mean      39.652498
std       13.436029
min       17.000000
25%       29.000000
50%       38.000000
75%       49.000000
max       90.000000
Name: age, dtype: float64

```

8. Among whom is the proportion of those who earn a lot (>50K) greater: married or single men? Among married those who have a marital-status starting with Married (Married-civ-spouse, Married-divorced, etc.) rest are considered bachelors.

```

data.loc[(data['sex'] == 'Male') &
         (data['marital-status'].isin(['Never-married',
                                       'Separated',
                                       'Divorced',
                                       'Widowed']))], 'salary'].value_counts()

```

```
↳ <=50K    7552
   >50K     697
   Name: salary, dtype: int64
```

```
data.loc[(data['sex'] == 'Male') &
         (data['marital-status'].str.startswith('Married')), 'salary'].value_counts()
```

```
↳ <=50K    7576
   >50K     5965
   Name: salary, dtype: int64
```

```
data['marital-status'].value_counts()
```

```
↳ Married-civ-spouse    14976
   Never-married        10683
   Divorced              4443
   Separated            1025
   Widowed              993
   Married-spouse-absent  418
   Married-AF-spouse      23
   Name: marital-status, dtype: int64
```

9. What is the maximum number of hours a person works per week (hours-per-week feature)? How many hours, and what is the percentage of those who earn a lot (>50K) among them?

```
max_load = data['hours-per-week'].max()
print("Max time - {0} hours./week.".format(max_load))
```

```
num_workaholics = data[data['hours-per-week'] == max_load].shape[0]
print("Total number of such hard workers {0}".format(num_workaholics))
```

```
rich_share = float(data[(data['hours-per-week'] == max_load)
                        & (data['salary'] == '>50K')].shape[0]) / num_workaholics
print("Percentage of rich among them {0}%".format(int(100 * rich_share)))
```

```
↳ Max time - 99 hours./week.
   Total number of such hard workers 85
   Percentage of rich among them 29%
```

10. Count the average time of work (hours-per-week) for those who earn a little and a lot (salary) these be for Japan?

```
pd.crosstab(data['native-country'], data['salary'],
            values=data['hours-per-week'], aggfunc=np.mean).T
```

```
↳
```

native-country	?	Cambodia	Canada	China	Columbia	Cuba	Dominican-Republic
salary							
<=50K	40.164760	41.416667	37.914634	37.381818	38.684211	37.985714	42.338235
>50K	45.547945	40.000000	45.641026	38.900000	50.000000	42.440000	47.000000

▼ Часть 2

```
!pip install -U pandasql
import pandas as pd
import pandasql as ps
from datetime import datetime
import seaborn
```

📄 Collecting pandasql

```
Downloading https://files.pythonhosted.org/packages/6b/c4/ee4096ffa2eeeca0c749b26f0
Requirement already satisfied, skipping upgrade: numpy in /usr/local/lib/python3.6/di
Requirement already satisfied, skipping upgrade: pandas in /usr/local/lib/python3.6/d
Requirement already satisfied, skipping upgrade: sqlalchemy in /usr/local/lib/python3
Requirement already satisfied, skipping upgrade: pytz>=2017.2 in /usr/local/lib/pytho
Requirement already satisfied, skipping upgrade: python-dateutil>=2.6.1 in /usr/local
Requirement already satisfied, skipping upgrade: six>=1.5 in /usr/local/lib/python3.6
Building wheels for collected packages: pandasql
Building wheel for pandasql (setup.py) ... done
Created wheel for pandasql: filename=pandasql-0.7.3-cp36-none-any.whl size=26819 sh
Stored in directory: /root/.cache/pip/wheels/53/6c/18/b87a2e5fa8a82e9c026311de56210
Successfully built pandasql
Installing collected packages: pandasql
Successfully installed pandasql-0.7.3
```

```
#Загружаем 2 датасета
from google.colab import files
files.upload()
```

📄

Выбрать файлы user_device.csv

```
• user_device.csv(application/vnd.ms-excel) - 9437 bytes, last modified: 06.04.2020 - 100% done
Saving user_device.csv to user_device.csv
{'user_device.csv': b'use_id,user_id,platform,platform_version,device,use_type_id\n22
```

```
from google.colab import files
files.upload()
```

📄

Выбрать файлы

user_usage.csv

- **user_usage.csv**(application/vnd.ms-excel) - 6432 bytes, last modified: 06.04.2020 - 100% done
Saving user_usage.csv to user_usage.csv
{'user_usage.csv': b'outgoing_mins_per_month,outgoing_sms_per_month,monthly_mb,use_id

```
data_1 = pd.read_csv('user_device.csv')
```

```
data_2 = pd.read_csv('user_usage.csv')
```

```
data_1.head()
```

	use_id	user_id	platform	platform_version	device	use_type_id
0	22782	26980	ios	10.2	iPhone7,2	2
1	22783	29628	android	6.0	Nexus 5	3
2	22784	28473	android	5.1	SM-G903F	1
3	22785	15200	ios	10.2	iPhone7,2	3
4	22786	28239	android	6.0	ONE E1003	1

```
data_2.head()
```

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id
0	21.97	4.82	1557.33	22787
1	1710.08	136.88	7267.55	22788
2	1710.08	136.88	7267.55	22789
3	94.46	35.17	519.12	22790
4	71.59	79.26	1557.33	22792

```
#pandas code
```

```
start_time = time.time()
```

```
join_df = pd.merge(data_1,
                    data_2[data_1.use_type_id == 1],
                    how = 'inner',
                    right_on = 'use_id',
                    left_on = 'use_id')
```

```
join_df = join_df[['use_id', 'user_id', 'platform', 'device', 'outgoing_mins_per_month', '
print(join_df)
```

```
print("--- %s seconds ---" % (time.time() - start_time))
```

```
↳
```

	use_id	user_id	platform	device	outgoing_mins_per_month	monthly_mb
0	22789	28714	android	SM-G930F	1710.08	7267.55
1	22792	28217	android	SM-G361F	71.59	1557.33
2	22793	28217	android	SM-G361F	71.59	1557.33
3	22794	28217	android	SM-G361F	71.59	519.12
4	22795	28217	android	SM-G361F	71.59	519.12
..
88	23041	28953	android	SM-G900F	198.59	5191.12
89	23043	28953	android	SM-G900F	198.59	5191.12
90	23044	28953	android	SM-G900F	198.59	3114.67
91	23046	29454	android	Moto G (4)	106.65	5191.12
92	23049	29725	android	SM-G900F	344.53	519.12

[93 rows x 6 columns]

--- 0.021392345428466797 seconds ---

#pandaSQL code

start_time = time.time()

join_query = ''

SELECT

d1.use_id,

user_id,

platform,

device,

outgoing_mins_per_month,

monthly_mb

FROM data_1 as d1 JOIN data_2 as d2 ON (d1.use_id = d2.use_id)

WHERE (use_type_id = 1)

...

pandasqlcode = ps.sqldf(join_query)

print(pandasqlcode)

print("--- %s seconds ---" % (time.time() - start_time))



	use_id	user_id	platform	device \
0	22787	12921	android	GT-I9505
1	22788	28714	android	SM-G930F
2	22789	28714	android	SM-G930F
3	22790	29592	android	D2303
4	22792	28217	android	SM-G361F
..
152	23043	28953	android	SM-G900F
153	23044	28953	android	SM-G900F
154	23046	29454	android	Moto G (4)
155	23049	29725	android	SM-G900F
156	23053	20257	android	Vodafone Smart ultra 6

	outgoing_mins_per_month	monthly_mb
0	21.97	1557.33
1	1710.08	7267.55
2	1710.08	7267.55
3	94.46	519.12
4	71.59	1557.33
..
152	198.59	5191.12
153	198.59	3114.67
154	106.65	5191.12
155	344.53	519.12
156	42.75	5191.12

[157 rows x 6 columns]

--- 0.034714698791503906 seconds ---

Как можно заметить код PandaSQL выполняется на 50% дольше

Группировка

#Код Pandas

```
start_time = time.time()
pandascode = pd.DataFrame(join_df.groupby('device').monthly_mb.mean())
pandascode = pandascode.sort_values('monthly_mb')
print(pandascode)
print("--- %s seconds ---" % (time.time() - start_time))
```



device	monthly_mb
HUAWEI CUN-L01	11.680000
HTC Desire 620	74.400000
GT-I8190N	407.010000
D2303	519.120000
HTC Desire 626	519.120000
SM-J320FN	778.665000
GT-I9300	894.580000
SM-G361F	934.404000
HTC One S	1038.210000
GT-I9195	1211.260000
iPhone7,2	1271.390000
SM-G935F	1384.303333
SM-A310F	1557.330000
SM-A500FU	1557.330000
SM-G903F	1557.330000
SM-G800F	1557.330000
SM-G920F	1557.330000
VF-795	1557.330000
SM-G360F	1557.330000
LG-H815	1557.330000
Lenovo K51c78	1557.330000
HTC Desire 530	1557.330000
C6603	1557.330000
D5503	1557.330000
EVA-L09	1557.330000
SM-G531F	2076.450000
ONE A2003	2076.450000
SM-A300FU	2336.000000
HUAWEI VNS-L31	3114.670000
ONEPLUS A3003	3823.610000
GT-I9505	3924.252857
SM-G925F	4152.880000
HTC Desire 825	4513.560000
SM-G900F	4556.647778
Moto G (4)	5191.120000
SM-N910F	5437.576667
D6603	7267.550000
SM-G930F	8305.775000
HTC One mini 2	10382.210000
X11	12458.670000
A0001	15573.330000
HTC Desire 510	15573.330000
GT-N7100	20764.450000
---	0.010747194290161133 seconds ---

#Код PandaSQL

```
start_time = time.time()
print(ps.sqldf('select "device", avg("monthly_mb") as AVG_Traf from pandasqlcode group
print("--- %s seconds ---" % (time.time() - start_time))
```



	device	AVG_Traf
0	HUAWEI CUN-L01	11.680000
1	HTC Desire 620	74.400000
2	MotoE2(4G-LTE)	212.640000
3	GT-I8190N	407.010000
4	GT-I9300	464.185000
5	D2303	519.120000
6	HTC Desire 626	519.120000
7	GT-I9506	803.240000
8	SM-J320FN	830.574000
9	SM-G361F	934.404000
10	HTC One S	1038.210000
11	GT-I9195	1211.260000
12	C6603	1557.330000
13	D5503	1557.330000
14	D5803	1557.330000
15	EVA-L09	1557.330000
16	GT-I9515	1557.330000
17	HTC Desire 530	1557.330000
18	LG-H815	1557.330000
19	Lenovo K51c78	1557.330000
20	Nexus 5X	1557.330000
21	SM-A310F	1557.330000
22	SM-A500FU	1557.330000
23	SM-G360F	1557.330000
24	SM-G800F	1557.330000
25	SM-G903F	1557.330000
26	VF-795	1557.330000
27	SM-A300FU	1687.112500
28	SM-G920F	1985.168000
29	F3111	2076.450000
30	ONE A2003	2076.450000
31	SM-G531F	2076.450000
32	HTC One M9	2362.070000
33	HUAWEI VNS-L31	3114.670000
34	SM-G925F	3633.775000
35	ONEPLUS A3003	3823.610000
36	SM-G900F	3841.427333
37	SM-G935F	4568.182000
38	E6653	5191.120000
39	Moto G (4)	5191.120000
40	Vodafone Smart ultra 6	5191.120000
41	HTC Desire 825	5498.970000
42	GT-I9505	5564.726364
43	HTC One_M8	6577.120000
44	D6603	7267.550000
45	SM-G930F	7959.700000
46	SM-N910F	8038.370000
47	GT-N7100	11939.560000
48	X11	12458.670000
49	HTC Desire 510	12562.488000
50	HTC One mini 2	13842.956667
51	A0001	15573.330000
52	SM-N9005	16611.550000
---	0.025293588638305664	seconds ---

Как можно заметить код PandaSQL выполняется в 2.5 раза больше, чем код Pandas